

Technical Overview on Micro-Services Platforms and Tools

Prakash Badhe
prakash.badhe@vishwasoft.in

About Me: Prakash Badhe ²

- ❑ Technology Practices Trainer, Mentor and Consultant.
- ❑ Technology Experience for 20+ years
- ❑ Passion for technologies, tools and frameworks
- ❑ Proficient in application frameworks, libraries and tools.
- ❑ Proficient in DevOps Tools and Technologies.
- ❑ Proficient in Docker, Kubernetes etc.
- ❑ Supports Agile technical practices and setup for agile development and production environments.
- ❑ Worked with teams from Symantec, Samsung, PTC, Cognizant, Oracle, Persistent, HP, Wipro etc.

Agenda

3

- Technical and Theoretical Overview on
 - MicroServices
 - MicroServices Environment and Tools
 - Docker Container Engine
 - Kubernetes for clustering containers
 - MicroServices in Kubernetes
 - Kubernetes in Cloud
 - Service Mesh Istio.
 - Continuous Integration and Continuous Delivery management with Jenkins
 - StackDriver
- Introduction to DevOps and DevSecOps

Software applications

4

- Desktop applications
 - Calculator
 - MS Excel
 - MineSweeper
- Mobile apps
 - Gmail
 - Samsung apps
 - ----
 - *Single User applications running locally.*
 - *Can access data from remote locations*
 - *Controlled by user.*
 - *Designed for local processing.*

Internet/Web applications

5

- Google , Microsoft , Oracle web site etc.
- www.google.com
 - Running 24 X 7 on Hosting Server.
 - Accessed by browser and other UI client applications.
 - Controlled by server admin.
 - Client-Server architecture
 - Multiple users can access the same simultaneously.
 - Works in standardized manner.
 - Standard language, Protocol, Network, Data standard etc.
 - Accessible anywhere in the world where internet network is available.
 - Standard protocol and data formats need to be defined.

Web Programming Model

- Involves client-server interaction on web.
- Uses http protocol and html for information exchange
- Exchange Messages that carry MIME-typed data
- Web applications are loosely coupled than the traditional distributed programming models like RPC, DCOM, and CORBA.

Web applications Limitations ⁷

- Accessible to only certain type of client applications like web browsers.
 - Large processing and storage capacity and maintenance needed on server side.
 - Need to be secured from unauthorized users and data corruption and leakage.
- Designed for Business to user and Business to Business for sharing the information via User interface.
- User interacts with the web site and fetch data or push data to server.
 - Performance is an issue in exchanging information
 - Not designed for application to application interactions.
 - Depends on specific platform for clients to access it

Application as Service

SOA : Service Oriented Architecture

Why SOA..?

- Distributed Processing..
- Applications need data sharing and communication in the network across different machines.
- Heterogeneous applications with different programming platforms and different Operating systems.
- Lots of data compatibility issues

Applications as services

- At abstract level if the applications are defined as platform independent services, then they can share the data and understand the data exchange...
- The SOA is the answer..
- Service Oriented Architecture.

Service applications

- Any client can access the service at any time from anywhere irrespective of platforms.
- Any application can access the service and share/exchange data.
- There is no hard coupling between applications with respect to application platform, OS, network etc.
- What matters is data/info exchange!
- Examples
 - Uber/Ola
 - Gmail on Mobile
 - E-commerce site backend data.
 - Logistics services

Service Oriented Architecture ¹²

- Client-server architecture.
- In a service-oriented architecture, applications are made up from loosely coupled software services, which interact to provide all the functionality needed by the application.
- The application services exchange information across the platforms.
- Example— any app can access the Google maps service.
- Each service is generally designed to be a self-contained and stateless to simplify the communication that takes place between them.

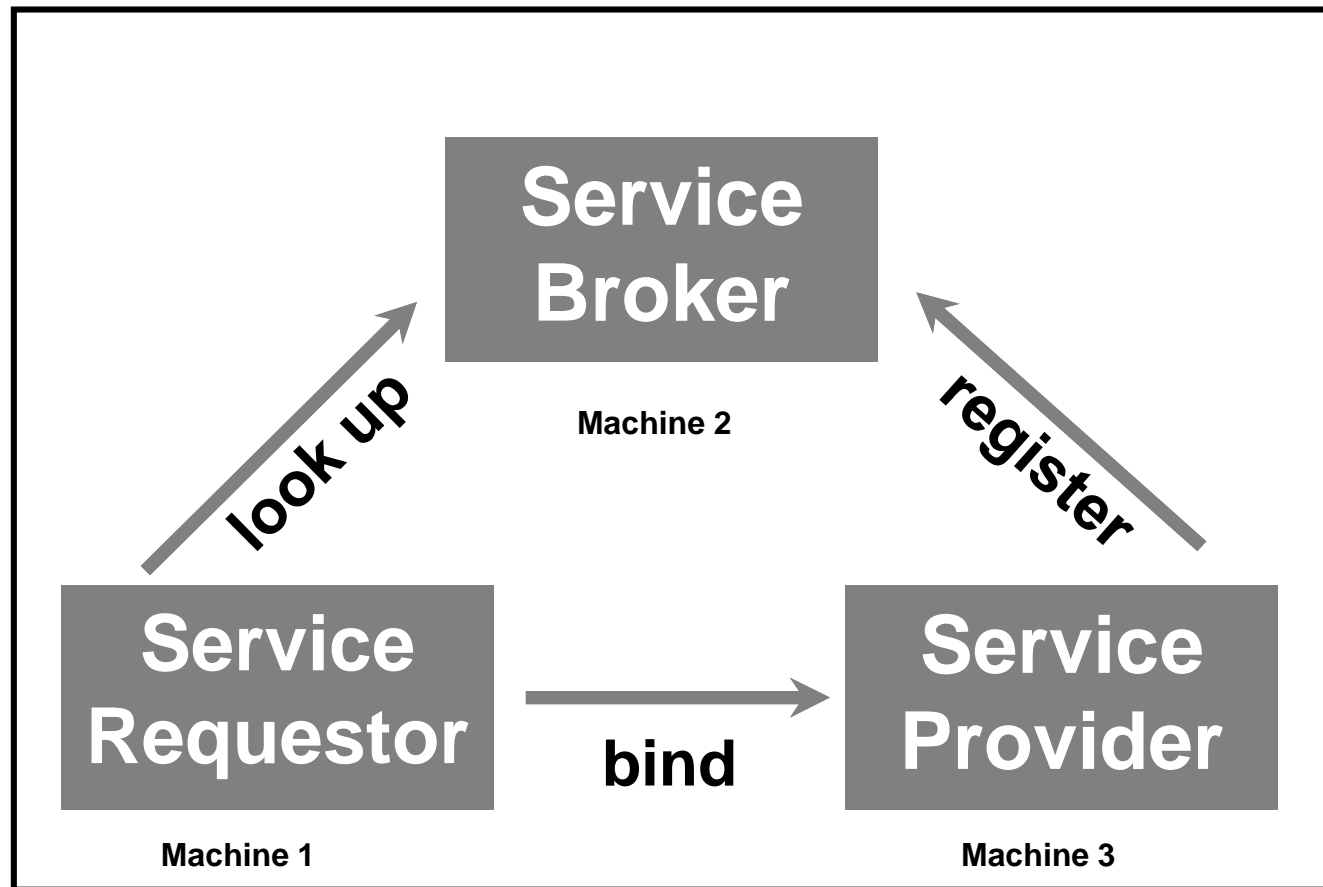
Roles Involved In A SOA

Three main roles involved in a service-oriented architecture

- **Service provider**
- **Service broker**
- **Service requestor**

SOA Communication Model

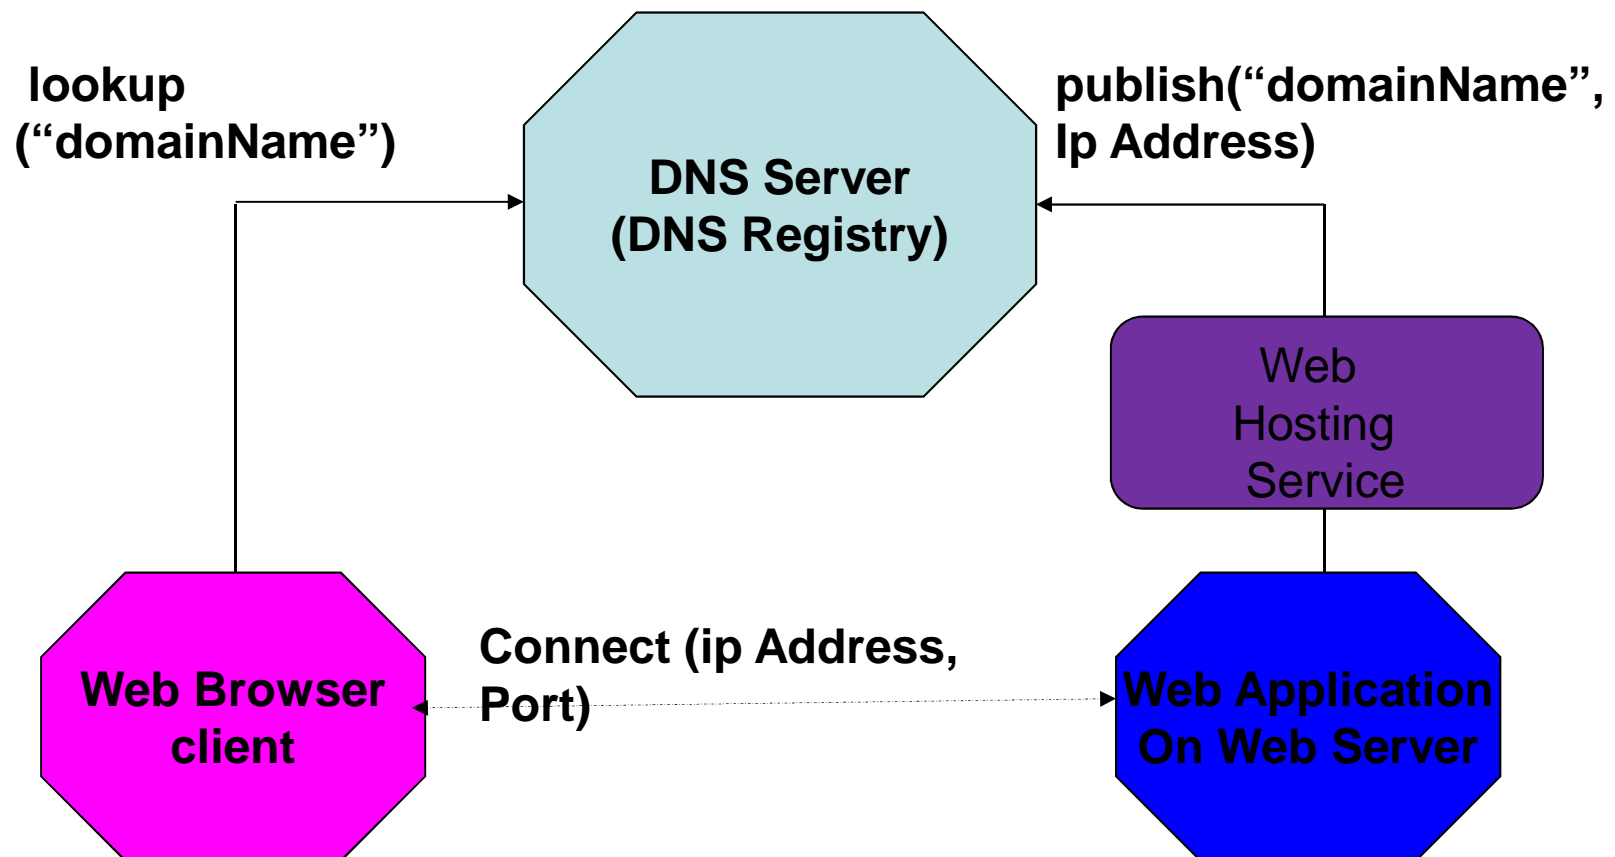
14



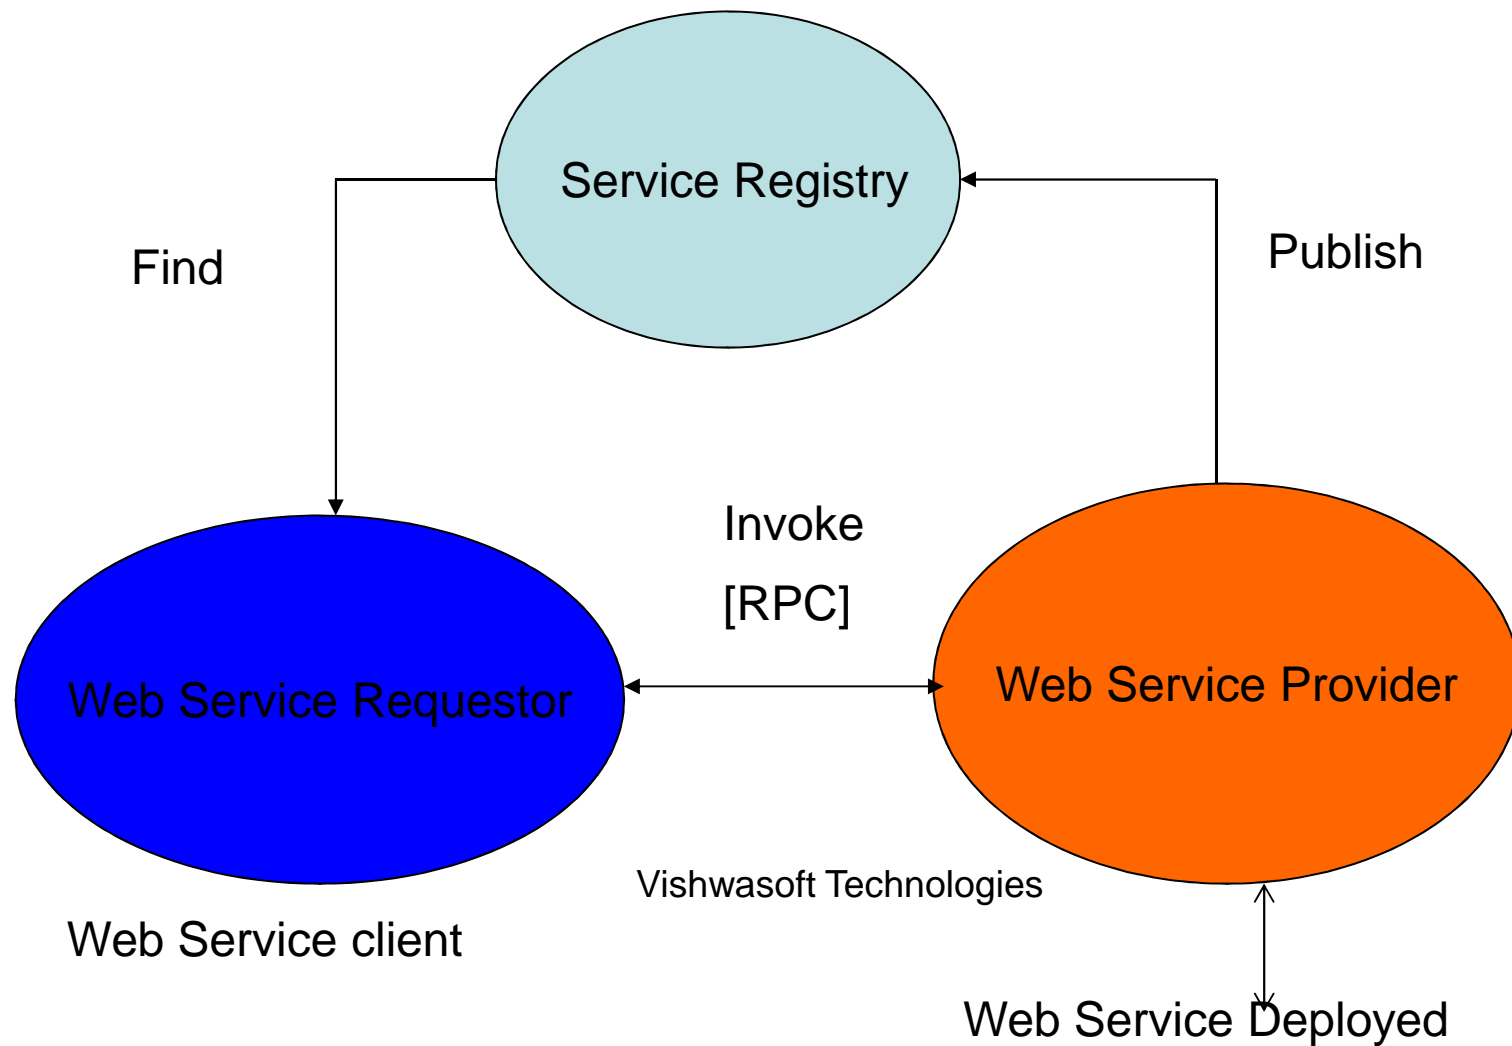
Service Oriented Architecture

Vishwasoft Technologies

Web Application Process



SOA Services Conceptually





Web Service As SOA

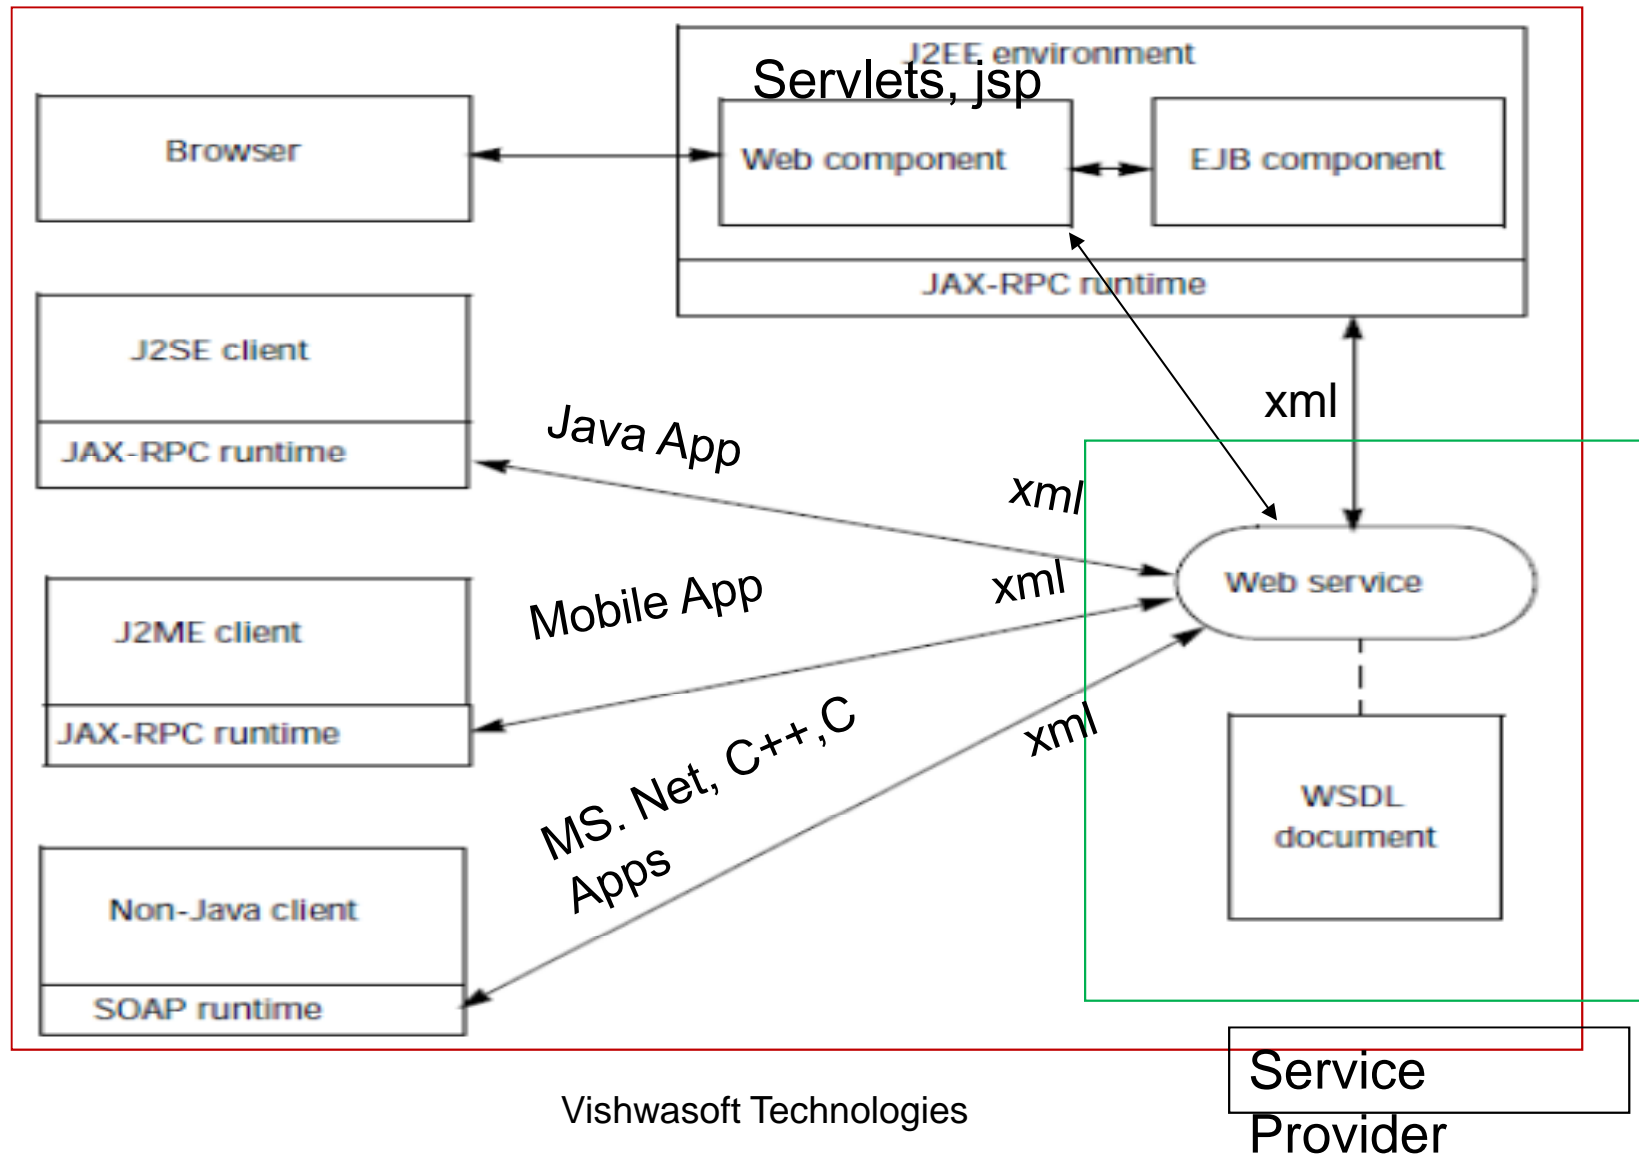
- A Web service is a software application designed to support **interoperable** machine-to-machine interaction over a network
- *Generally web services are small task /service operations supported by applications on the server/s.*
- *These tasks are reusable across different OS platforms and programming language applications.*
- *Web services operate in a distributed environment.*

Web service features

- Loosely coupled Web programming model for use in applications that are **not browser-based**.
- Provides a platform for building distributed applications using software components that are..
 - running on different operating systems and devices,
 - written using different programming languages and tools from multiple vendors,
 - all developed and deployed independently

Client Applications

20



XML Usage

- Web Services fundamentally use XML to standardize the behaviors and data.
- Key XML standards to understand
 - XML : programming language independent and extensible mark up across programming languages.
 - XML Schema specifies xml structure.
 - XML Namespace avoids naming conflicts and categorizes elements as per the functionality

JSON : New Data Format

22

- Text only Lightweight data format
- Data in arrays/maps format
- Quick to transport
- Easy/quick to parse
- Relational data supported.

Web Service Advantages

- Manage language independent interactions across different platforms and applications
- Different versions evolved over the time to support additional features.
- Interoperability is still an issue across different implementations and versions.
- Heavy dependencies on the soap library and other xml implementations.
- Non-xml data format support is limited only in the form of attachments.
- Performance is an issue for mission critical applications.

Web Service Standards

- SOAP
 - – Simple Object Access Protocol for web service and clients to communicate.
- WSDL
 - – Web Service Definition Language for web services to describe their services and other information.
- UDDI
 - – Universal Description, Discovery, and Integration protocol for Web Services to discover web services.
- All these are specified in XML format.

S

RESTing the SOAP?



Enter 'REST' Services

- REST is 'Representational State Transfer'..
- REST is an architectural style rather than a protocol which removes the dependencies on soap and xml standards as wsdl, uddi etc.
- REST supports any understandable data formats across applications.
- REST works currently only with Http.
- REST specifies Resources/data on server rather than actions on them..
- REST specifies transfer of the state of Resource across applications.

Resources on the Web

27

Resources are not only represented as XML

- XML formats (HTML, XHTML, RSS, etc.)
- JPG, GIF, PNG
- MP3, WAV, OGG
- Anything else that can be on the web.

Resource Nouns

- Important ‘things’ (nouns) are Resources
 - Addressed through a URI
- Uniform interface (verbs)
 - In HTTP: GET, PUT, POST, DELETE
- Verb-noun séparation makes intégration easier
 - GET /customer/45 Instead of getCustomer(45)

REST Nouns and Verbs

- REST works with resources as nouns with their identities.
- The state of these nouns is shared with clients over http methods as verbs(operations)
- The http methods as verbs in REST
 - Get: get the resource state/values
 - Post : post new resource
 - Put : update the resource state
 - Delete :delete the resource on server
 - Options: read the options available with server.
 - Head : set the http headers on server.

The **CRUD** operations denoted by http methods used from clients.

REST Commandments

- Give every “thing” an ID
- Link things together
- Use standard methods
- Communicate statelessly

SOAP and REST

- REST is ready for the enterprise
- REST is strong at:
 - Internet scale computing
 - High levels of interoperability
 - Resource Oriented operations
- SOAP/WS is strong at:
 - Complex security (Trust and Federation)
 - Multi-transport services
 - Occasionally connected applications
- In the real world they are typically enabled by a combination of Soap and REST

REST Frameworks

- Jax-RS is the java specification standard for implementing and consuming REST web services.
- Implementations are Jboss RestEasy, Jersey platforms.
- Spring with Spring-Boot
- Oracle and IBM SOA Suite
- Eclipse Micro-Profile
- MS.Net
- PHP/Python

SOA Applications

- Applications as services.
- Applications developed with different programming languages and running on different platforms able to communicate each other and exchange information/data.
- Applications as Services collaborate together as separate modules and share data.
- SOAP web services
- REST services
- Shared data as xml, json, text etc.

Multilayered Applications

- Presentation — responsible for handling HTTP requests and responding with either HTML or JSON/XML (for web services APIs).
- Business logic — the application's business logic.
- Database access — data access objects responsible for access the database.
- Application integration — integration with other services (e.g. via messaging or REST API).

Monolithic applications

Monolithic application is a single-tiered software application where the user interface and data access code are combined into a single program running on a single platform as single process.

A monolithic application is self-contained and independent From other computing applications.

Monolithic application is designed without modularity.

Monolithic application is not scalable and difficult to maintain And develop across teams.

Single Monolith - All in One

36

- Sometimes even with the modular layered components, the application is packaged and deployed as a single monolith. (All in One)

Monolith Benefits

- Simple to develop.
- Simple to test. For example you can implement end-to-end testing by simply launching the application and testing it.
- Simple to deploy. You just have to copy the packaged application to a server.
- Simple to scale horizontally by running multiple copies behind a load balancer.

Monolith Drawbacks

- Limitation in size and complexity.
- Application is too large and complex to fully understand and to be able to update with new changes fast and correctly.
- The size of the application can **slow down the start-up** time.
- The **response time and performance** becomes slow because of all execution happening in same process.
- The entire application has to be re-deployed on each update and during down time nothing of it can be accessible.

Monolith drawbacks..

- Difficult to develop across teams.
- Difficult to scale dynamically.
- Difficult to upgrade to new requirements, features.
- Difficult for Continuous deployments
- Difficult to increase the scope
- Run time performance issues
- Crashes the entire application due to small bugs.
- Porting/migration issues

Monolith - Reliability

Bug in any module (e.g. memory leak) can potentially bring down the entire process.

Moreover, since all instances of the application are identical, that bugs impact the availability of the entire application.

Scaling the Monolith

41

Challenge in scaling when different modules have conflicting resource requirements.

Monolith- Sharing and Reuse

42

- The good parts of one application cannot be shared or re-used by other applications.
- Sharing of database for other processes and parallel updates is a challenge.


Break the Monolith...

- The application architecture is broken into a set of smaller, interconnected services (isolated processes) called as micro-services.
- Some micro-services expose a REST, RPC or message-based API and most services consume APIs provided by other services.
- Some micro-services might implement a web UI.
- The multiple modules work as service applications and collaborate together to form the application backbone.

Monolith


OR

Microservices




Micro Service –Service Layers ⁴⁵

Monolithic Architecture

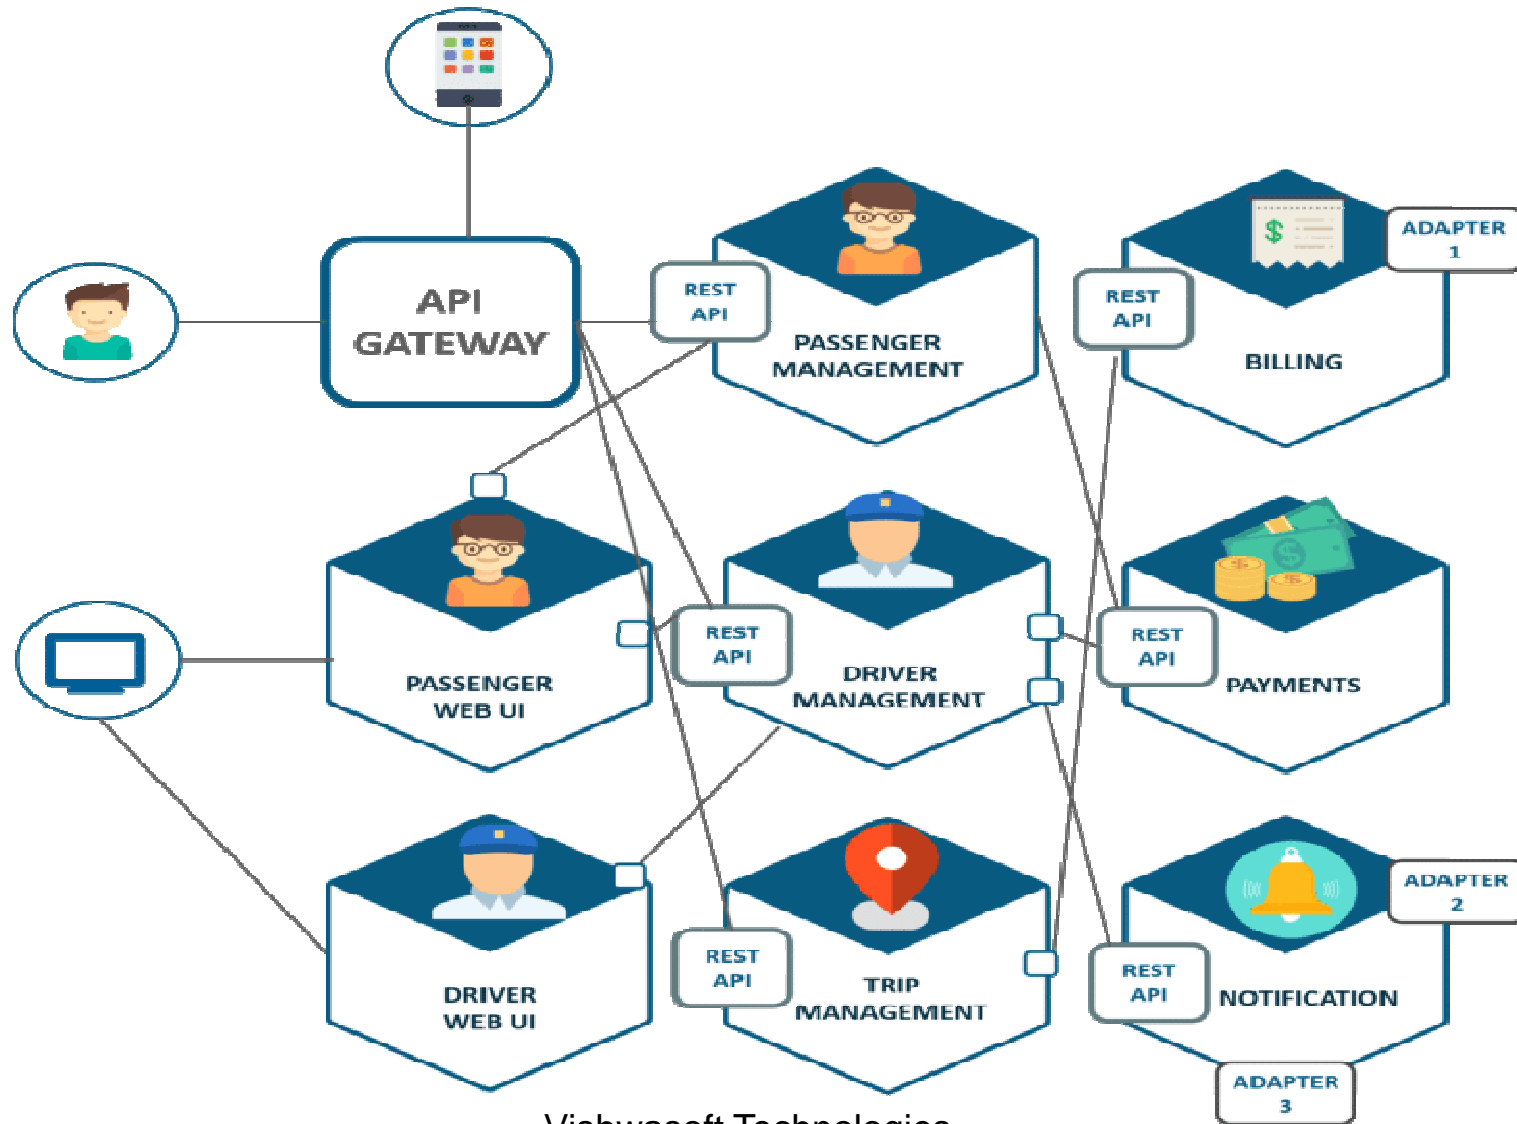


Microservices Architecture



MicroService applications

46




Monolith to Micro-Services

Monolithic



Microservices

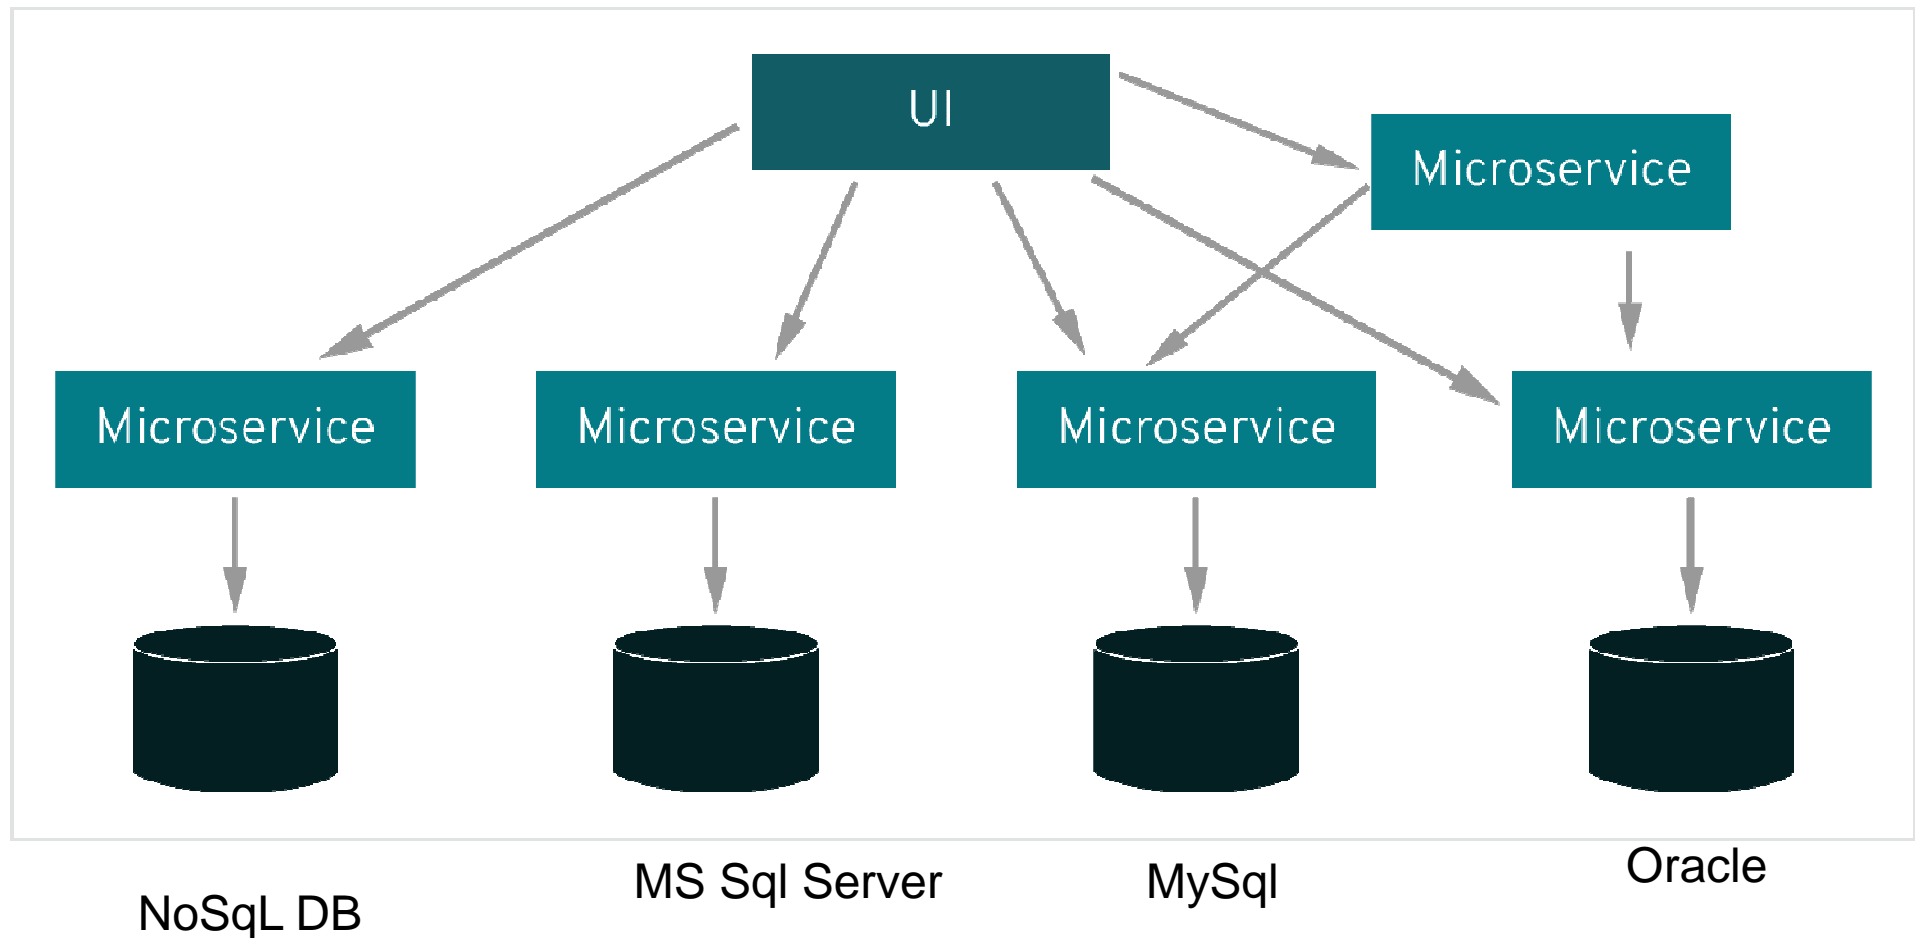


Monolith to MicroServices

- MicroService is an architecture pattern for service applications to collaborate together .
- Suitable to develop across team of developers.
- Easier to develop with smaller modules and smaller teams as well.
- Highly testable because of loose coupling with other services/modules
- Easier to scale
- Easier to troubleshoot
- Flexible in code maintenance.
- Independently deployable units.

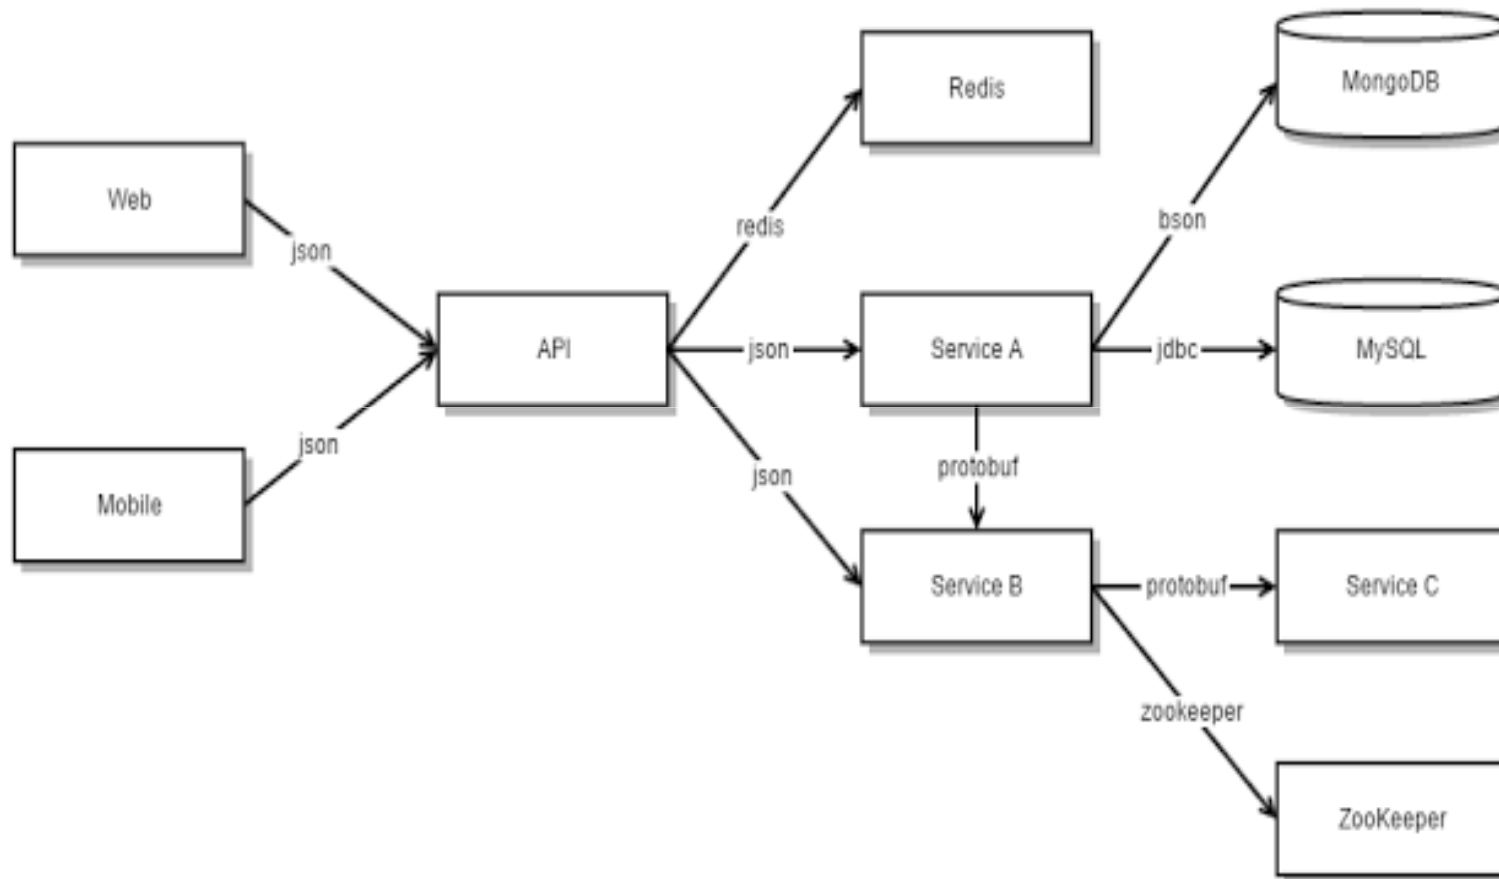
MicroService applications

49



Service layers in Micro Services

50



MicroService Design Principles⁵¹

- ✓ Loose coupling
- ✓ Scalability
- ✓ Availability
- ✓ Resiliency: ability to recover quickly from difficulties
- ✓ Flexibility
- ✓ Failure isolation
- ✓ Independent, autonomous
- ✓ Decentralized
- ✓ Failure isolation
- ✓ Auto-Provisioning of Nodes
- ✓ Continuous Delivery


Decompose by...

- Decompose by business capability and define independent services corresponding to business capabilities.
- Decompose by domain-driven design paradigm with isolated model design.
- Decompose by verb/action or use case and define separate services that are responsible for particular actions.
- Decompose by nouns/resources by defining separate a service that is responsible for all operations on particular entities/resources of a given type.e.g. Account, Users etc.
- Every service should have only a small/single set of responsibilities

How to de-compose


53

Identify Modules by domain, categories, business use, related dependencies.




Database break-up

Break out database tables, wrap with service and update dependencies




Service Modules

55



UI Extraction

56



MicroServices in application

57

- The architecture that structures the application as a set of loosely coupled, collaborating services.
- Highly maintainable and testable
- Loosely coupled with other services
- Independently deployable
- Each service is developed by a small isolated team.

MicroService Platforms

- Spring with Spring Boot (Development and deployment)
- Spring Boot with Spring Cloud and Netflix(Deployment)
- Cloud Foundry supports multi-cloud application platform as a service (PaaS) governed by the Cloud Foundry Foundation.
- Spring Cloud with Cloud-foundry
- Eclipse MicroProfile.(development and deployment)
- Microsoft .NET platform for API(Develop and deploy)
- AWS Cloud , Google CE, Azure, OpenShift (RedHat) with Kubernetes as deployment platforms.
- Lagom MicroService platform

MicroService Benefits

- MicroServices reinforce modular structure, which is particularly important for larger teams.
- Services are easier to deploy, and since they are autonomous, are less likely to cause entire system failures when they go wrong.
- With MicroServices you can mix multiple technologies, languages, development frameworks and data-storage technologies.
- Reduced complexity.
- Each MicroService is to be deployed independently. As a result, it makes continuous deployment possible for complex applications

MicroService Benefits

- Performance is improved, since it is running as multiple isolated processes instead of one like in monolith.
- Better testability — services are smaller and faster to test.
- Better deploy-ability — services can be deployed independently.
- Improved fault isolation; for memory leak in one service then only that service is affected.
- MicroService architecture enables each service to be scaled independently.

MicroService Challenges

61

- Manage the complexity of creating a distributed system.
- Manage the inter-service communication mechanism and deal with failures.
- Implementing requests spanning across multiple services is more difficult.
- Testing the services interactions is tedious.
- Troubleshooting at High level with integration issues .
- Increased resource consumption as memory , CPU since each service is a separate isolated process.

MicroService Patterns

- Patterns are style, specifications for the MicroServices on how to collaborate and exchange the data.
- Patterns are implemented across the different platforms and programming environments.
- The infrastructure supports the services and features required for MicroServices.

Integration Patterns

63

- ❖ **Aggregator Pattern**
- ❖ **Splitter**
- ❖ **UI separation**
- ❖ **Monitoring patterns**
 - ❖ **Log Aggregation**
 - ❖ **Performance Metrics for analysis**
 - ❖ **Distributed Tracing**
 - ❖ **Health Check**
- ❖ **Command Query Responsibility Segregation(CQRS)**
- ❖ **Asynchronous processing with Message Brokers.**
- ❖ **Event sourcing**

Deployment Patterns

64

- ❖ Externalized Configuration
- ❖ Dynamic Service Discovery
- ❖ API Gateway
- ❖ Circuit Breaker
- ❖ Blue-Green Deployment
- ❖ Saga for consistency
- ❖ Load balancing
- ❖ High availability

Deployment Environments

- NetFlix
- Spring Cloud
- Google Cloud, AWS, Azure
- Moleculer Framework based on Node JS
- Microsoft .Net platform
- Python/Django
- PHP frameworks

Data management patterns

- Database per service
- Shared database across services
- Saga for transactions across services.
- Event sourcing
- Command Query Responsibility Segregation(CQRS)
- API Composition as aggregate from multiple services.

Database Per service

- To ensure loose coupling, each service can have its own isolated database.
- Instead of sharing a single database across the services, **each service has its own database.**
- **T**this can result in duplication of some data. However, having a **database schema per service is essential** to get the benefit from micro-services, because it **ensures loose coupling.**

Challenges for Database per service

- Main challenges include implementing business transactions that span several MicroServices.
- Another challenge could be implementing queries that want to expose data from two or three different bounded contexts across the services.

Maintain Data consistency

69

- Maintaining data consistency across services is a challenge since the 2 phase-commit/distributed transactions is not an option for many services.
- In certain Enterprise applications, the replication across the database instances is done in background to maintain consistency.
- In certain cases the services can opt for shared database across multiple services.

Shared Database

- This approach uses a shared database accessed by multiple MicroServices.
- The ACID transactions are used to enforce consistency.
- Developers across teams need to coordinate for schema changes to tables.
- The run-time conflicts/data corruption can occur when multiple services are trying to update the same database resources.

Cost of MicroServices

- Distributed systems are harder to implement, since remote calls are slow and are always at risk of failure. Uniform programming model can reduce the difficulties.
- Maintaining strong consistency is extremely difficult for a distributed system, which means everyone has to manage eventual consistency. With high availability nodes in clusters, this effect can be reduced.
- Need a mature operations team to manage lots of services, which are being redeployed regularly, with automation CI/CD this can be eased.
- The network, technology, bandwidth limits the performance and flexibility.

Infrastructure for MicroServices⁷²

- The environment for MicroServices deployment and execution.
- API Gateway for Request-Response routing and security authentication.
- Load balancing : dispatch the requests to available nodes by calculating the service load or in round robin manner.
- Dynamic service discovery of services independent of the location/ip address by implementing service registry like domain name service (DNS) for web applications.
- Should support service health monitoring API points.

Infrastructure support...

73

- The fault tolerance features like Circuit Breaker, Bulkhead, Retry, Timeout, Fallback etc.
- Middleware for integration, Cache etc.
- Transporters for node to node communication in a cluster.
- Should support documentation about the services and their external APIs and data formats etc.
- Data Serializers and de-Serializers with custom schemas.

Clients for MicroServices

74

- In a server-side enterprise application, It should support a variety of different clients including desktop browsers, mobile browsers and native mobile applications.
- The application can also expose an REST API for third parties to consume.
- It might also integrate with other applications via either web services or a message broker.

Process inside the MicroService⁷⁵

- There are logical components corresponding to different functional areas of the service
 - The service handles requests (HTTP requests and messages) by executing
 - The business logic process
 - Access a database
 - Exchange messages with other systems
 - Return a HTML/JSON/XML response.

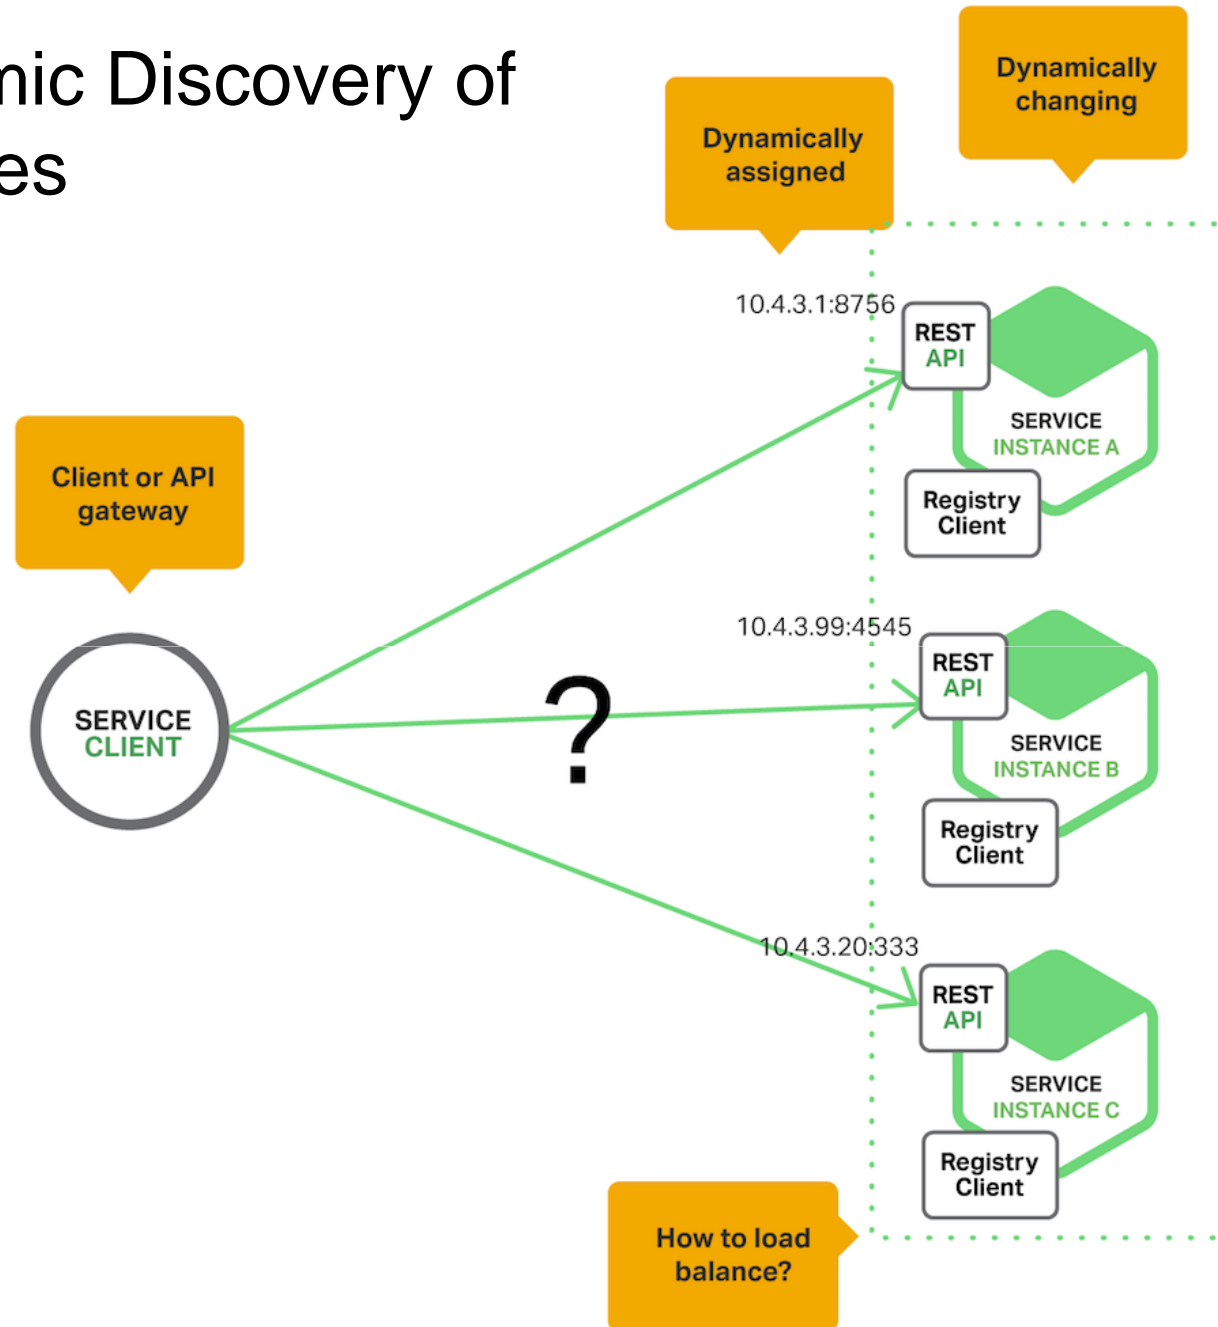
- .

Service Registry and services⁷⁶

- Services should be isolated and independent.
- Services should NOT have hard coded references to other services.
- **Services should be able to register and discover from external Service Registry environment.**


Dynamic Discovery of services

77




Self Registration

78




Third Party Registration

79



API Gateway

80




API Gateway Responsibilities ⁸¹

- Routing of requests to specific machine/node
- Security authentication and authorize for access control
- Load balance the requests across nodes/machines

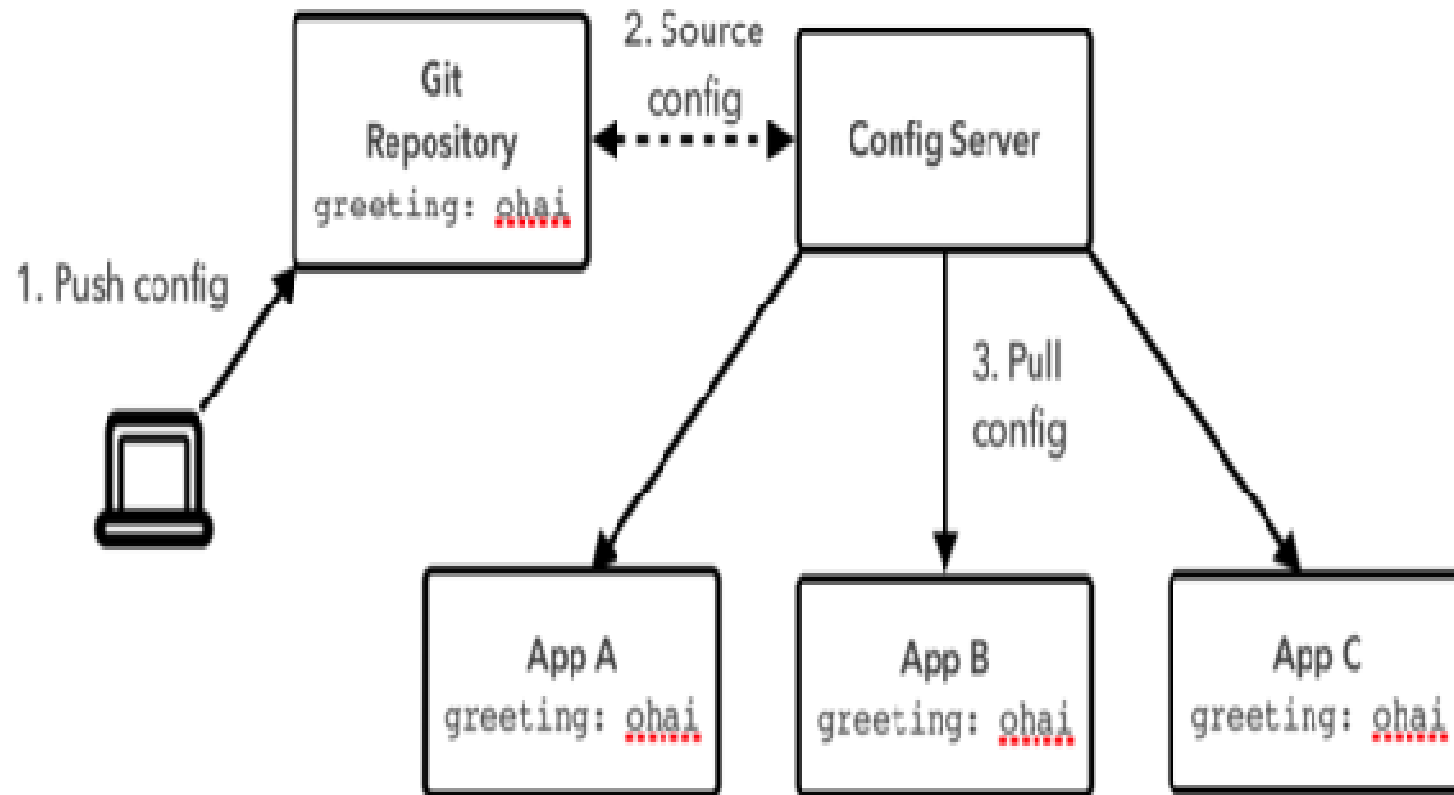
Externalized Configuration

82



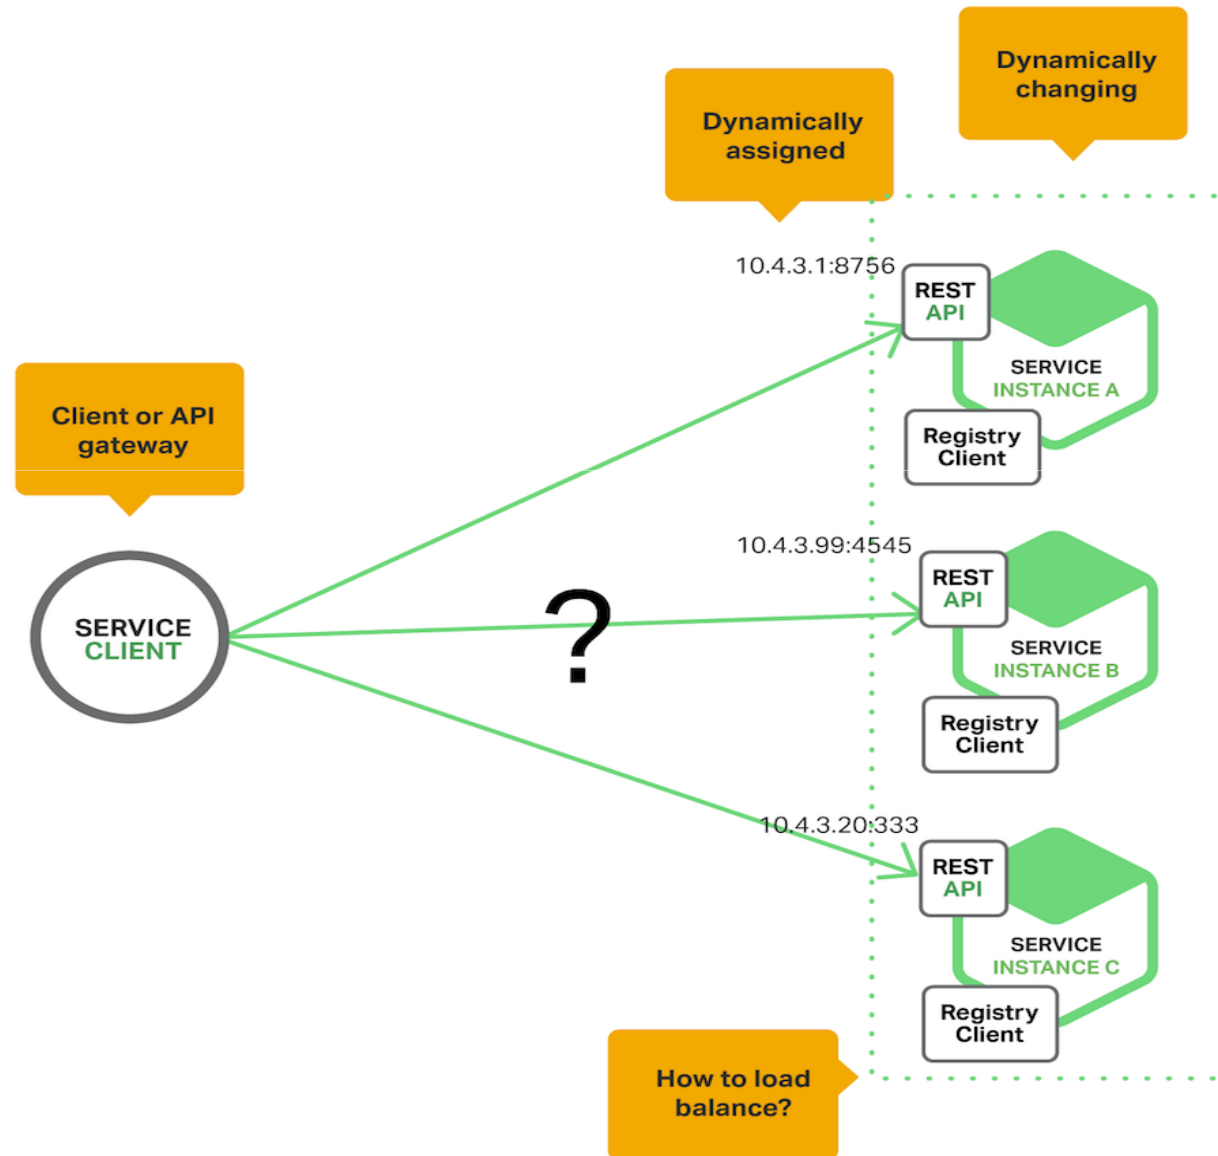
Getting the configuration

83



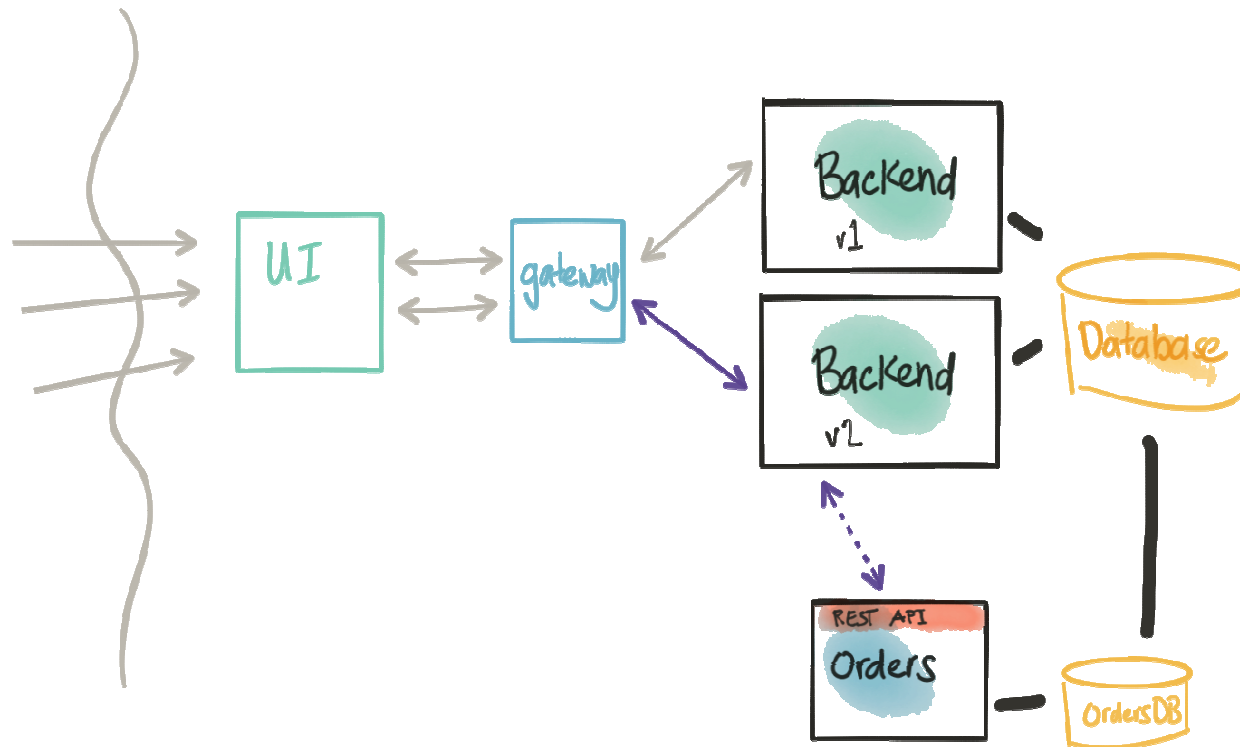
Balancing the Load... Across machines

84

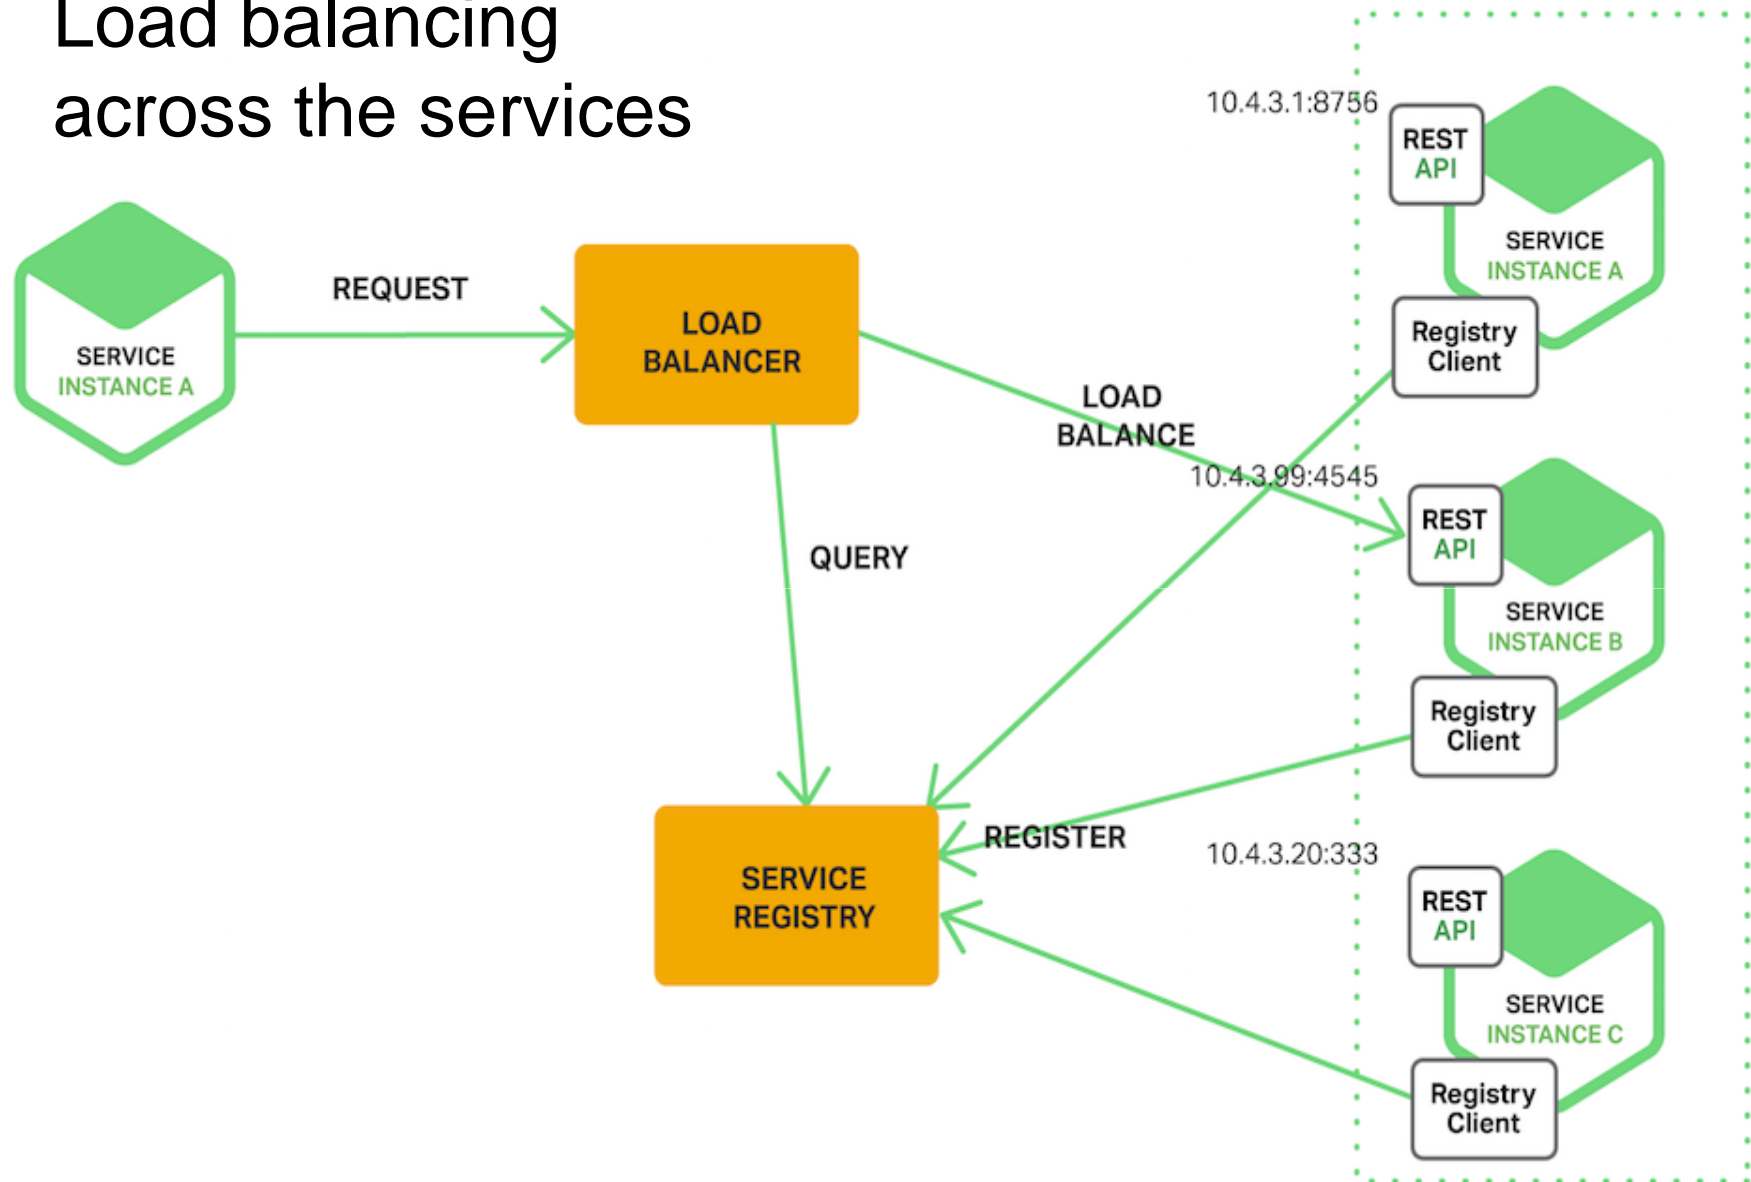


Load Balancing

85



Load balancing across the services



Algorithms for Load balancing ⁸⁷

- Dispatch across nodes in Round Robin manner.
- Weighted dispatch depending on weight of the node calculated based on capacity, request load etc.
- Sticky node etc.

Handle Error Gracefully

88

- When services synchronously invokes another there is always the possibility that the other service is unavailable or is responding with high delay.
- The resources such as threads might be consumed in the caller while waiting for the other service to respond and this leads to resource exhaustion, which makes the calling service unable to handle other requests.
- The failure of one service can cascade to other services throughout the application.
- How to prevent a network or service failure from cascading to other services?

Circuit Breaker

89

- Parent service invokes the remote service via a proxy that functions similar to an electrical circuit breaker.
 - When the number of consecutive failures crosses a threshold, the circuit breaker trips.
 - For the duration of a timeout period all attempts to invoke the remote service will fail immediately.
 - After the timeout expires the circuit breaker allows a limited number of test requests to pass through.
 - If those requests succeed, the circuit breaker resumes normal operation.
 - Otherwise if there is a failure the timeout period begins again.
- Parent service handles the failure of the services that they invoke.
- **Netflix Hystrix** is an implementation for this pattern.

Need of Atomic Updates

- Service to update the database **and** send messages/events to message broker/system.
- The database update and sending of the message must be atomic in order to avoid data inconsistencies and bugs. E.g if the database update fails, message update should roll back or so.
- One of the option is a distributed transaction (like JTA) that spans the database and the message broker to atomically update together, but may be time consuming and cause performance bottlenecks.

Transactions across services ⁹¹

- Majorly database per service applications.
- Saga pattern is the solution to implement business transactions spanning multiple MicroServices.

SAGA for distributed transactions ⁹²

- A **Saga** is a sequence of local transactions.
- For every transaction performed within a Saga, the service performing the transaction publishes an event.
- The subsequent transaction is triggered based on the output of the previous transaction.

Chain the transactions..

- If one of the transactions in this chain fails, the Saga executes a series of compensating transactions to undo the impact of all the previous transactions.
- To understand this better, let's take a simple example.

Saga in Food-delivery app

94

- A customer tries to order food, the below process starts..
 - Food Order service creates an *order*. At this point, the order is in a PENDING state. A Saga manages the chain of events.
 - The Saga contacts the restaurant via the Restaurant service.
 - The Restaurant service attempts to place the order with the chosen restaurant. After getting confirmation, it sends back a reply.
 - The Saga receives the reply. And, depending on the reply, it can Approve the order or Reject the order.
 - The Food Order service then changes the state of the order. If the order was Approved, it would inform the customer with the next details. If Rejected, it will also inform the customer with an apology message.


Time lost in communication..

- The services collaborate to handle client requests.
- The services should use an inter-process communication mechanism and agreed data structure of information exchange.
- The synchronous communication between the services results in tight runtime coupling, both the client and service must be available for the duration of the request and creates performance bottlenecks.

Messaging Applications

- In enterprise applications there are different types of messaging systems and they need to interoperate and communicate each other.
- Most of the systems are proprietary which makes it difficult for integration and portability.
- J2ee defines an open messaging standard system which works seamlessly with other systems and defines new messaging systems and services.

Application Messaging in Enterprise systems



Message Broker

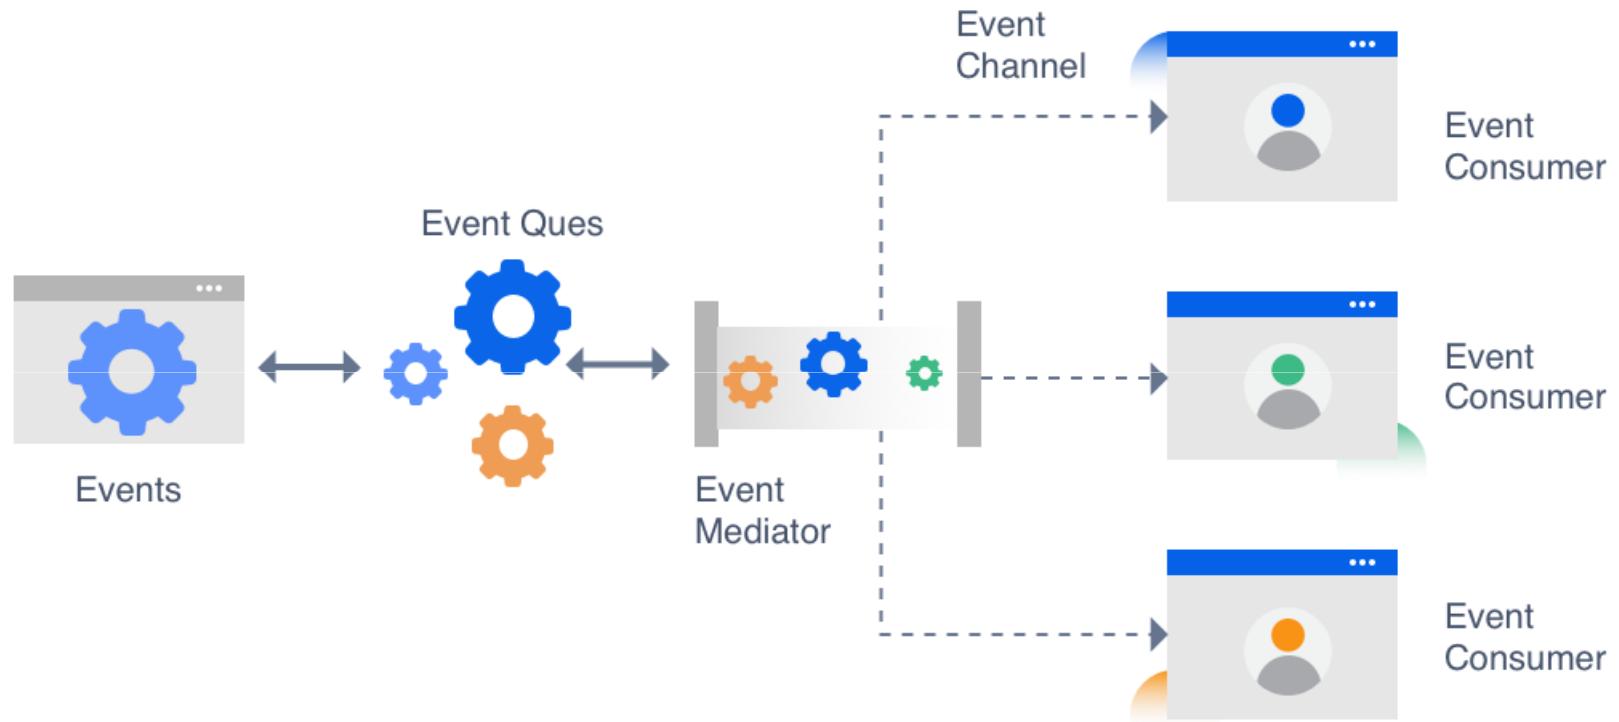
- The services use **asynchronous** messaging for inter-service communication.
- Services communicate by exchanging messages over messaging channels.
 - Request/response mode - a service sends a request message to a recipient and expects to receive a reply message promptly
 - Notifications - a sender sends a message a recipient but does not expect a reply. Nor is one sent.
 - Request/asynchronous response - a service sends a request message to a recipient and expects to receive a reply message eventually.

Messaging over channels

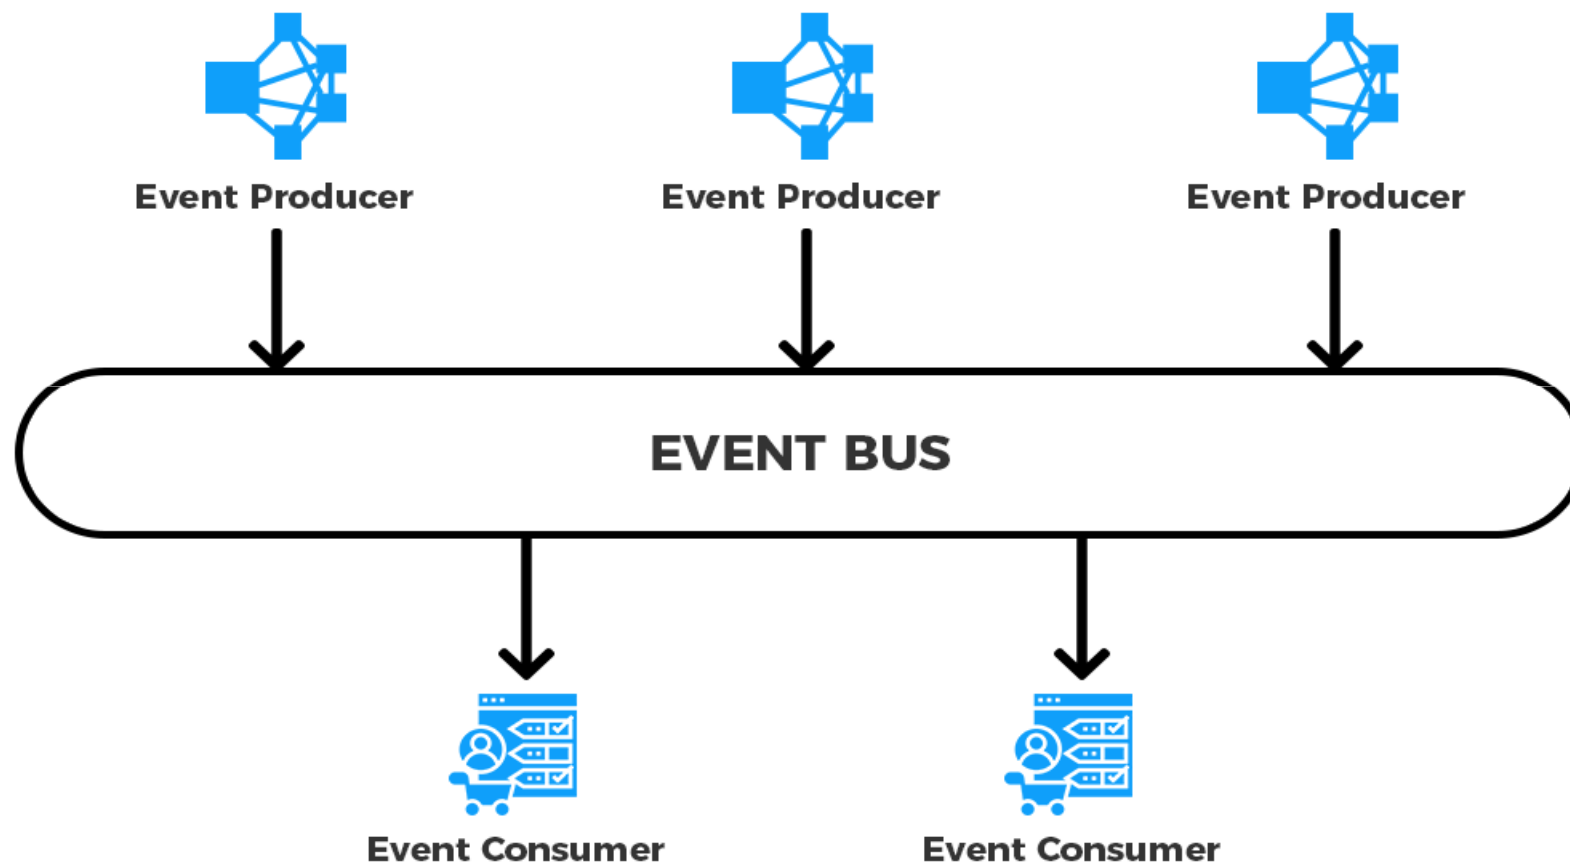
- Publish/subscribe - a service publishes a message to zero or more recipients through common message destinations
- Publish/asynchronous response - a service publishes a request to one or recipients, some of whom send back a reply.
- Asynchronous messaging brokers
 - Apache Kafka
 - RabbitMQ
 - Apache ActiveMQ
 - JMS Messaging

Event Driven Architecture

100



Event Bus

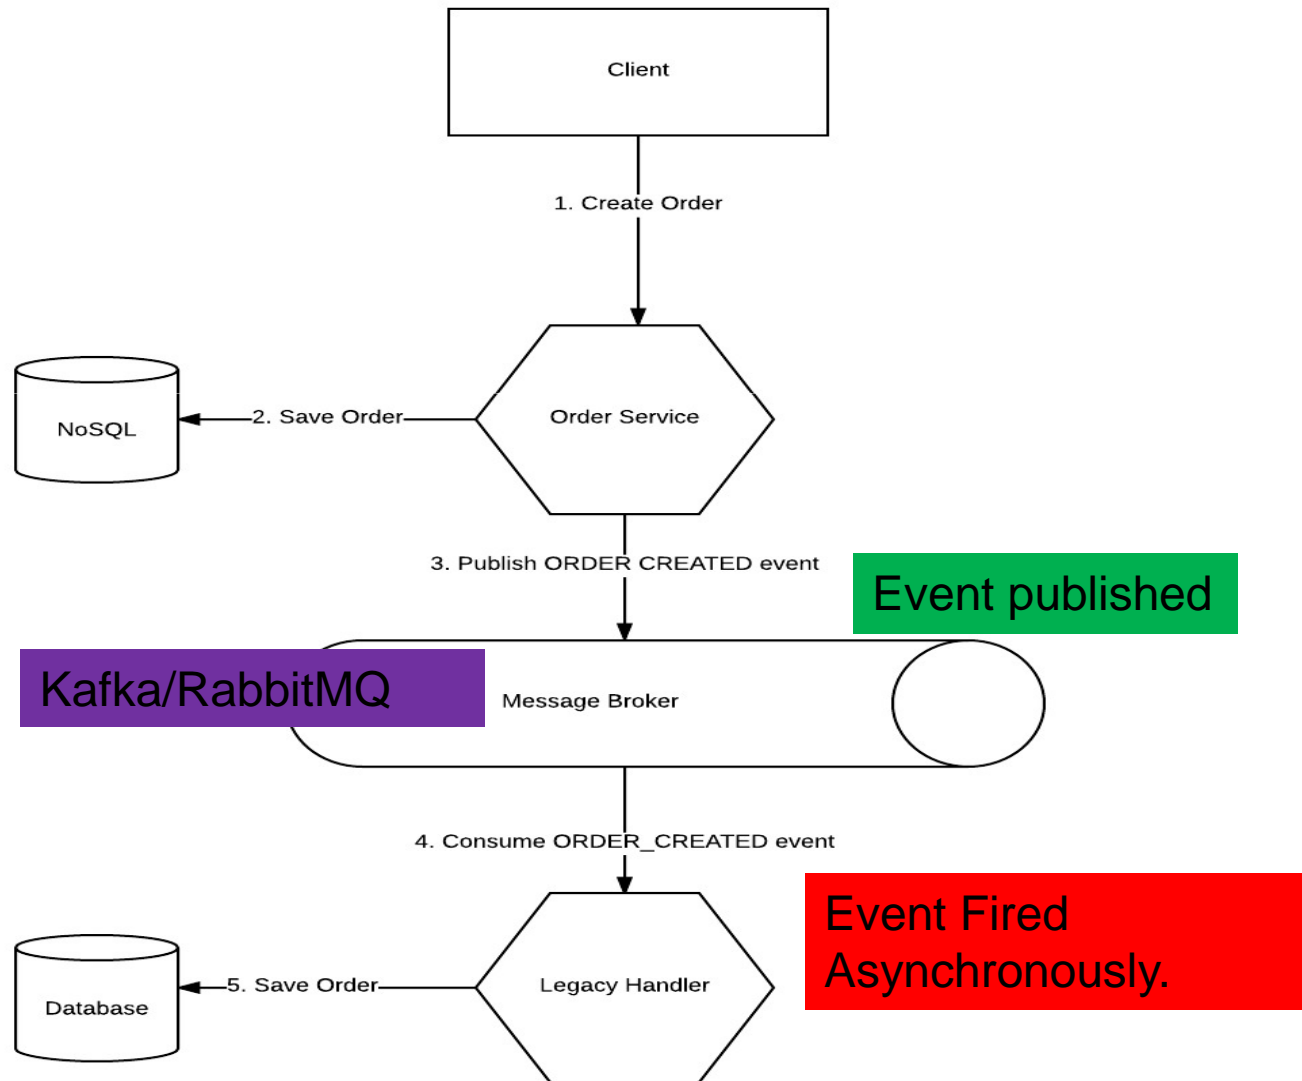


Event Driven Micro-Services

102

- Child service called by parent service or clients if it takes more time for execution, it blocks the clients or parents services.
- Apply event driven approach with messaging in Message Brokers like ActiveMQ, RabbitMQ, JMS servers or Apache Kafka.

Event Through Message Broker¹⁰³



Messaging benefits and Issues¹⁰⁴

- Loose runtime coupling between services since it decouples the message sender from the consumer.
- Improved availability since the message broker buffers messages until the consumer is able to process them
- Supports a variety of communication patterns including request/reply, notifications, request/async response, publish/subscribe, publish/async response etc.
- **This adds additional complexity to message broker, which must be highly available.**


Event Sourcing

105

- To reliably/atomically update the database and publish messages/events.
- Event sourcing saves the state of a business entity such as Order or a Customer as a sequence of state-changing events.
- Saving an event is a single and inherently atomic operation.
- The application can reconstruct an entity's current state by replaying the events.

Event Sourcing

106

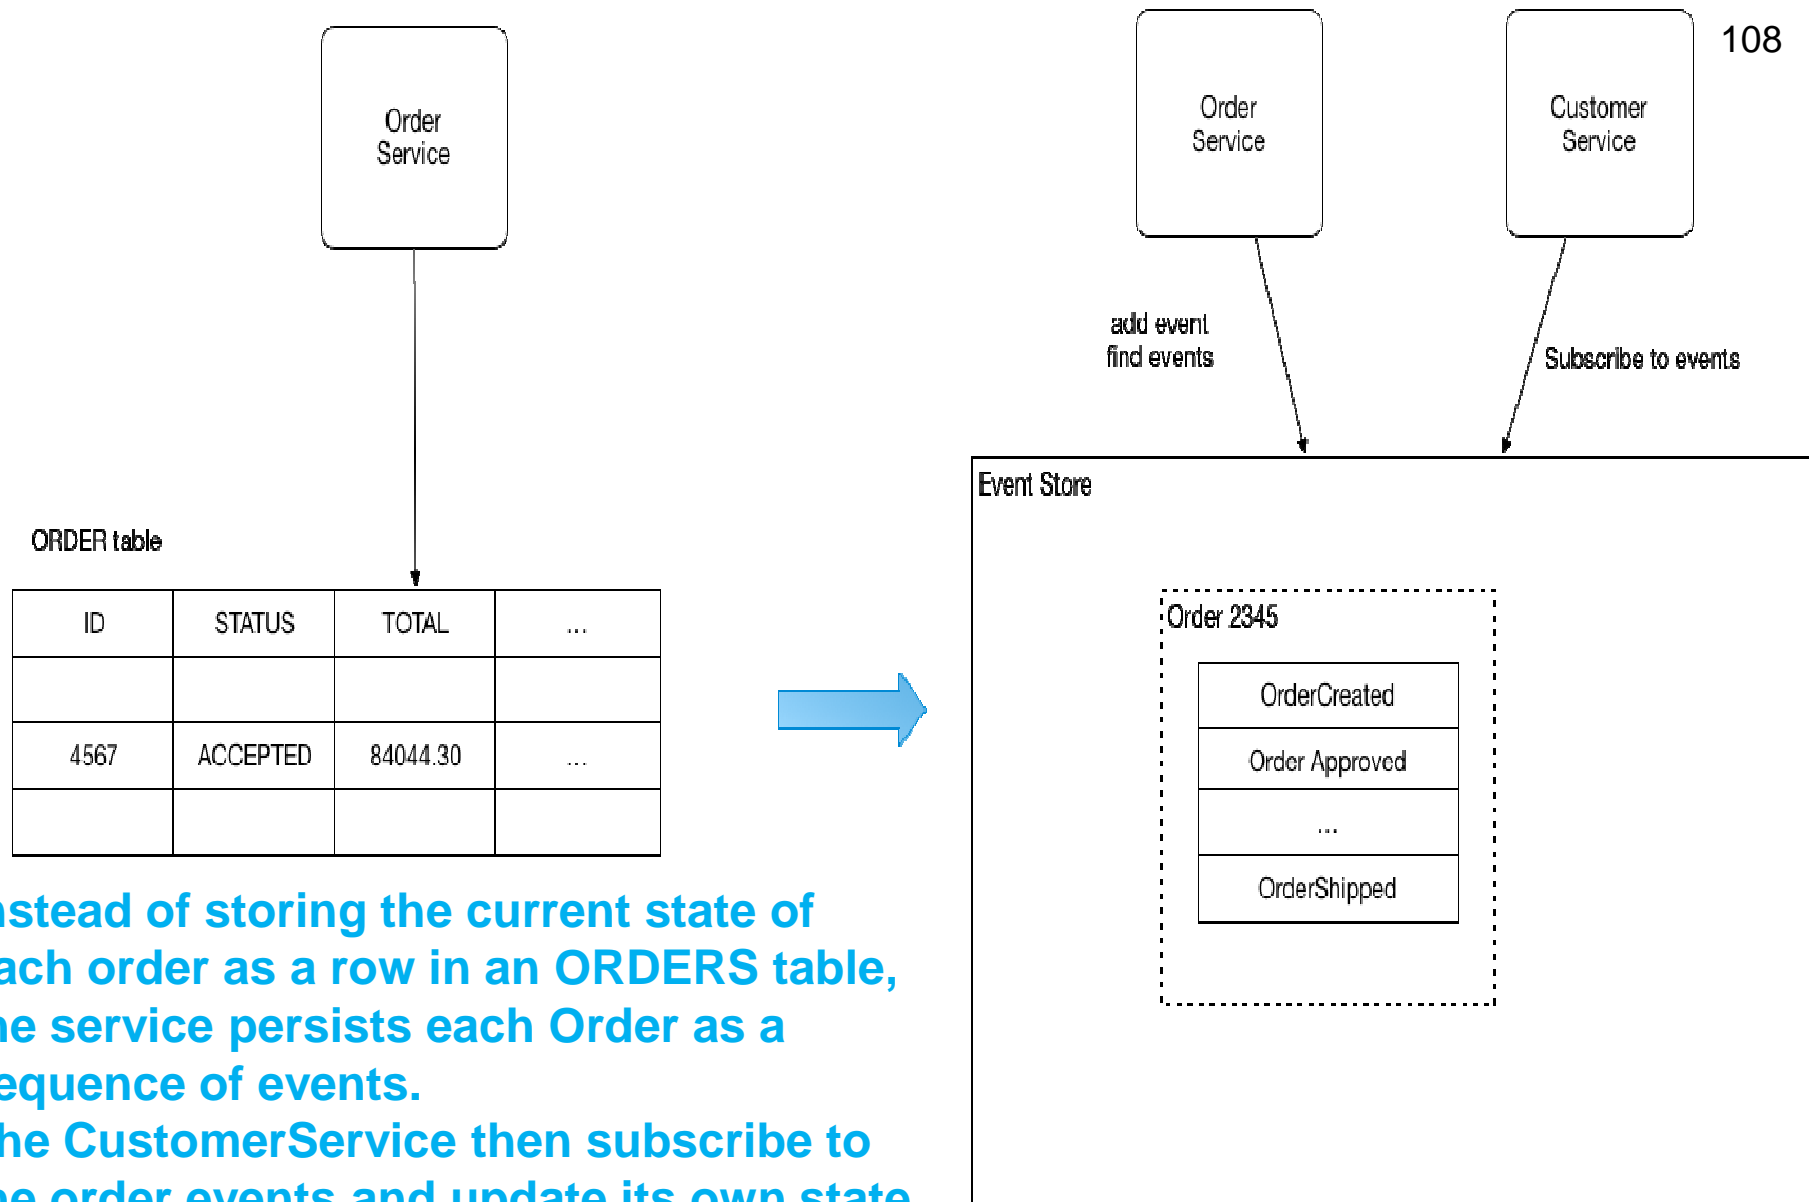


Event Sourcing in Action

107

- Persist events in an event store, which is a database of events.
- The store has an API for adding and retrieving an entity's events and the event store also behaves like a message broker.
- It provides an API that enables services to subscribe to events.
- When a service saves an event in the event store, it is delivered to all interested subscribers.

Event Sourcing with Message Broker



108

Instead of storing the current state of each order as a row in an ORDERS table, the service persists each Order as a sequence of events.
The CustomerService then subscribe to the order events and update its own state.

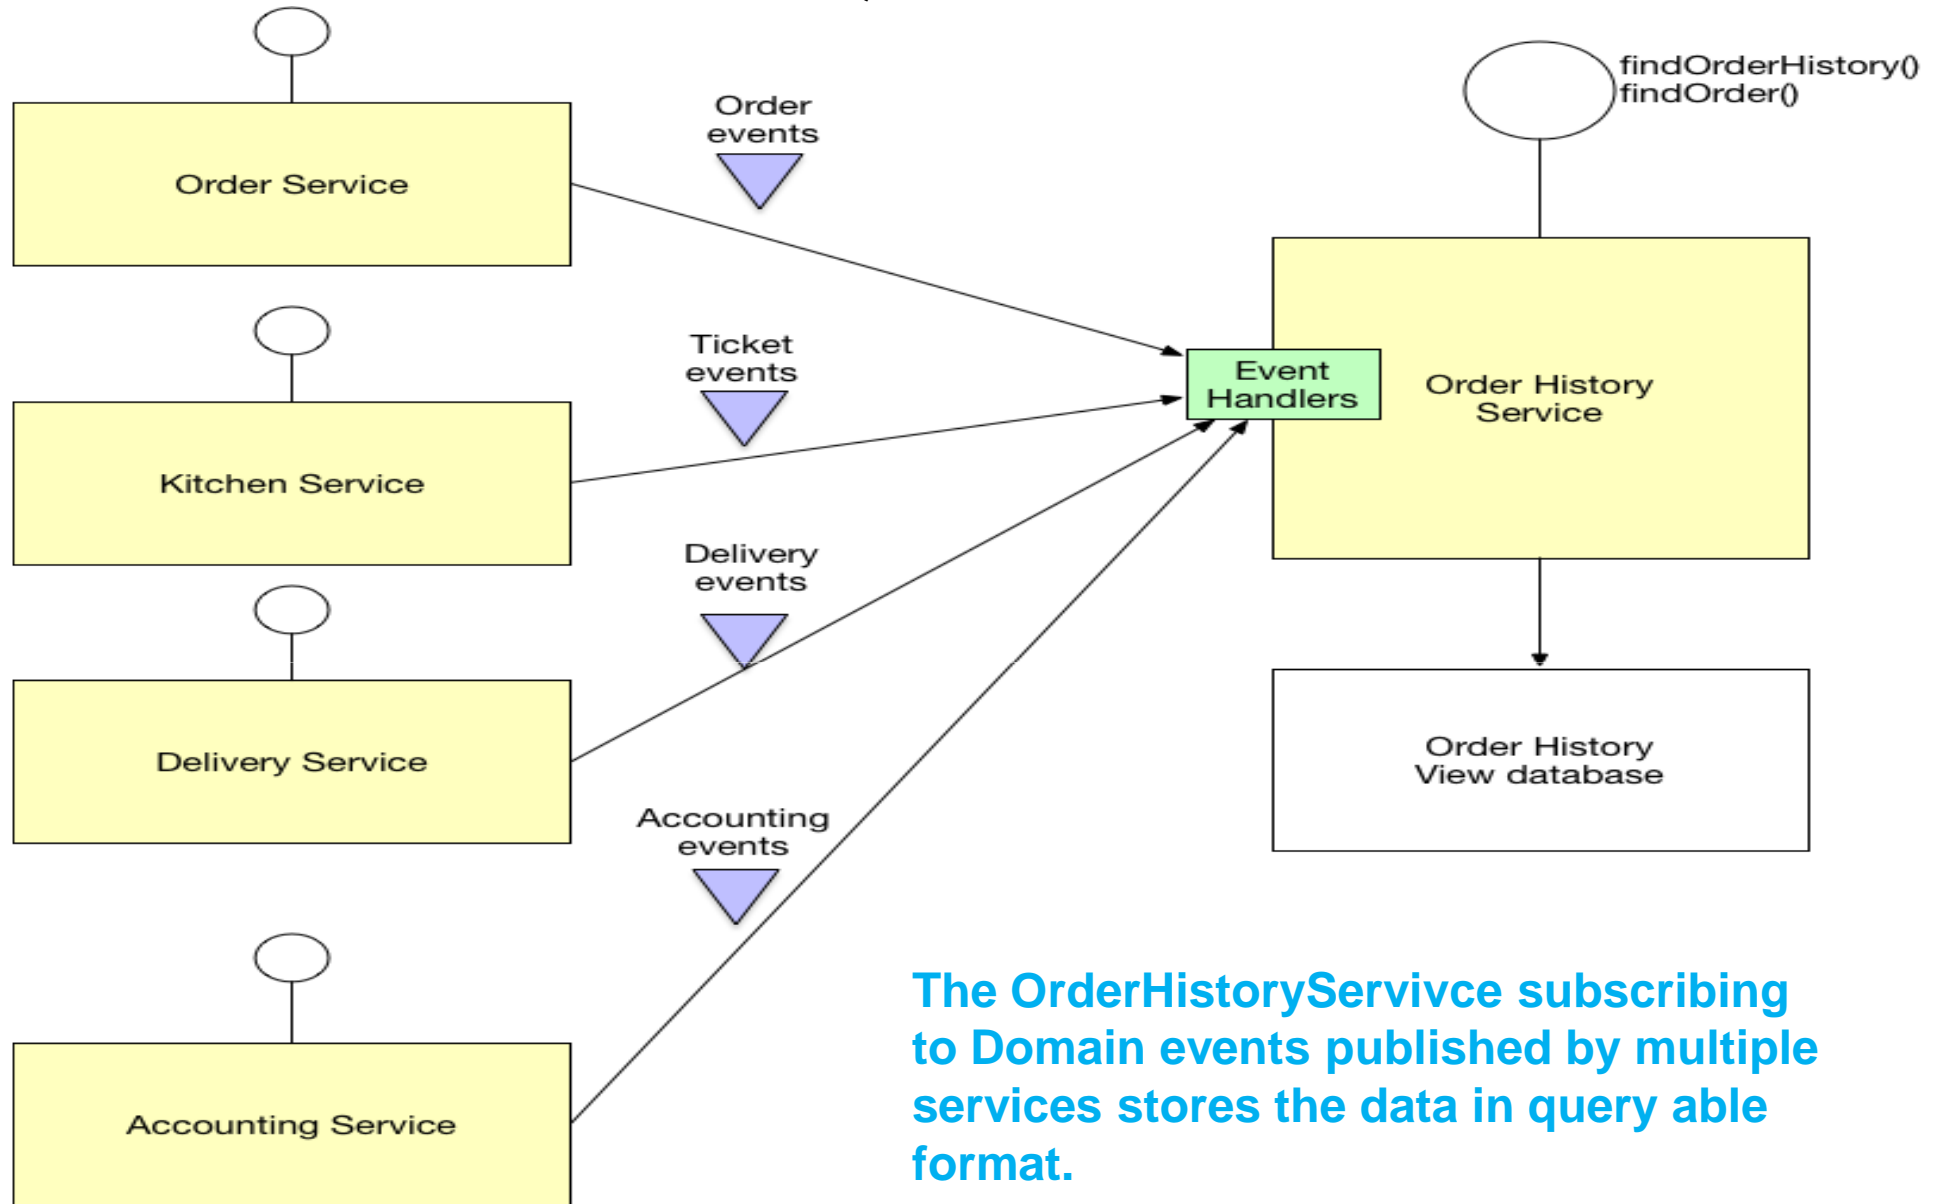
Query with the data

- To implement queries that joins data from multiple services.
- In case of Event sourcing pattern then the data is no longer easily query-able.
- A query system that enables to retrieve data from multiple services in a MicroService architecture.

Command Query Responsibility Segregation (CQRS)¹¹⁰

- Multiple services publish data as events to event source using event handlers.
- The event handlers are subscribed by single Event Source service.
- The view database - a read-only (for faster queries) supports the aggregate join queries.
- The Event Source service application keeps the data upto date by subscribing to Domain events published by the services that owns the data.

CQRS Pattern



Distributed Tracing

112

- Helps gather timing data needed to troubleshoot latency problems in service architectures.
- Data collection and analysis.
- Data can be structured.
- To quickly locate the data query based systems needed.
- Data can be summarized based on attributes.

Tracing with ZipKin

- Zipkin is a distributed tracing system.
- It helps gather timing data needed to troubleshoot latency problems in service architectures.
- Features include both the collection and lookup of this data.
- Data can be collected with trace ID in a log file, or jump directly to it.
- Run a query can query based on attributes such as service, operation name, tags and duration.
- Data can be summarized based on attributes.

Zipkin : Inputs for data

114

- The data can be reported to Zipkin via http or Kafka and many other options such as Apache ActiveMQ, gRPC and RabbitMQ.
- The data served to the UI is stored in-memory, or persistently with a supported backend such as Apache Cassandra or Elasticsearch.

ZipKin UI

115

- The Zipkin UI presents a Dependency diagram showing how many traced requests went through each application.
- This is helpful for identifying aggregate behavior including error paths or calls to deprecated services.

Spring Cloud Sleuth

- Spring Cloud Sleuth implements a distributed tracing solution for Spring Cloud,.
- Data can be captured in logs, or by sending it to a remote collector service.

Logging with ELK

- LogStash to collect the logs from different sources.
- The log collection can be load balanced.
- Passes the log data to Elasticsearch for storage.
- Elasticsearch supports analysis queries over RERST API.

All in One → Consul

118

- Consul is a service networking solution to connect and secure services across **any runtime platform and public or private clouds**.
- Supports service discovery, authorization, gateways as proxies and lot more as runtime platform.
- Consul supports service applications in platform independent ways.

Best Practices for MicroServices¹¹⁹

- Implement REST APIs in better optimized ways with patterns for loose coupling and performance and design.
- Utilize the swagger documentation or **HATEOAS** (navigable, linked restful API docs) for end clients to understand it.
- Utilize actuator like libraries for runtime analysis.
- Wherever possible add fallbacks with Circuit Breakers
-

Best Practices..

- Utilize dynamic discovery and distributed service configuration.
- Apply Load balancing.
- Think of reusability at design time from across platforms.
- Implement continuous delivery with Jenkins/Travis/Bamboo or TFS.
- Implement Monitoring and alerts with tools like Nagios at every possible API.
- Enable logging for critical points.
- Implement log segregations and analysis with ELK stack

Best Practices..

- Application performance management with Zipkin.
- Apply Event sourcing and CQRS at architecture level.
- Independently develop and deploy services.
- Keep service small.
- Utilize asynchronous workers for backend operations.

Best Practices

- Protect Resource APIs with security.
- Use lightweight data formats as JSON.
- Utilize NO-SQL Database for backend.
- For frequently updated services maintain versioned services.
- Apply limits on resource utilizations.
- Design stateless services for performance.

Happy Micro-Serving..